

 Welcome to Snorkel


Welcome to Snorkel!

We are happy to have you aboard. This document will guide you through the necessary steps to start contributing to our projects. Any additional questions [can be answered in our FAQ](#).

Platform Navigation & Functionality

Sign into our platform —>>>[Snorkel Expert Platform](#)<<<—

You will see this landing page, **sign in with the email address where you received your project invite**



First-time User Sign-In

Once you're added to your first project in the Snorkel Expert Platform, you will **receive an email with your login credentials and a temporary password**.

The temporary password will be a string of characters that may include punctuation marks at the end – You will need to include all letters, numbers, punctuation, and special characters in the password.

You will be asked to change your password once you login.

If you don't receive a temporary password contact the support team

Start collaborating

Once you are logged in you will see our home screen, here you will see three options per project.

At the start of every project you will only be enabled to do SUBMISSION task

The screenshot shows the Snorkel web interface. On the left is a sidebar with navigation links: Home, Marketplace, Dashboard, and an ADMIN section containing Dashboard, Projects, Users, Tasks, Metrics, EC Override, Packaging, Payments, and Assignment debug. The main content area is titled 'My projects' and shows '3 projects available'. A search bar contains 'copperhead' and a dropdown menu is set to 'Newest First'. Below this, three project task cards are listed for 'Project: Copperhead_Production': 'Review', 'Adjudication', and 'Submission'. Each card has a 'Start' button and an 'Unpublished' status tag. The 'Submission' card is highlighted with a red rectangular border.

Slack

We use Slack as our main communication tool, you can use slack on your browser or download the app

If this is your first time using Slack or working with Snorkel, you will receive an invitation to join the workspace, please check your email.




Accept the invite to Snorkel Expert Contributors

Once you're added to your first project in the Snorkel Expert Platform, you will **receive an email with your login credentials and a temporary password.**

The temporary password will be a string of characters that may include punctuation marks at the end – You will need to include all letters, numbers, punctuation, and special characters in the password.

You will be asked to change your password once you login.

1. Once you are assigned to the project you will be added to the to **#copperhead-ec-announcements** channel. On this channel we will share updates, feedback and general project announcements.



Project canvas & resources

Once added to the Slack channel it's important to navigate to our **Copperhead channel canvas**. There you will find everything to start you project onboarding

Slack Code of Conduct

It is important to the Snorkel team that our Slack channels remain a comfortable and welcoming environment for communication between Expert Contributors (ECs) and Snorkelers. The channels are intended to provide an accessible way for ECs to receive updates, ask questions, raise concerns, and collaborate with other experts and the Snorkel team.

To maintain that environment, we are dedicated to the following:

- We will maintain respectful and professional relations with Snorkel staff and other ECs.
- Any threats against Snorkel staff and/or other ECs will not be tolerated.
- Use of profane and/or obscene comments will not be tolerated.

The purpose of the Code of Conduct is not to prevent ECs from voicing dissatisfaction; if an EC feels something is wrong we encourage them to report it. We ask that ECs do so in a manner aligned with the ethos for this community & that promotes a productive dialogue amongst peers and/or Snorkel.

If violations of this code of conduct occur, we will address the initial violation with the EC privately.

If the violations persist, Snorkel will take action up to and including revoking an EC's Slack access or offboarding them from further Snorkel work. ECs will still be able to communicate with Snorkel staff via email.

Thank you all for continuing to participate in polite and constructive conversations on Slack. We truly appreciate your contributions to both the project and to our EC community!

 Welcome to Copperhead Reviews

What This Project Is About

The goal of this project is to build computer-use agent (CUA) environments and tasks that reflect real, complex professional work. These tasks represent the multi-step workflows people perform across desktop applications, terminal environments, and web services to produce a concrete deliverable.

Within the Snorkel contributor platform, you will be provided with everything needed to evaluate a task, including the task instructions, supporting files, expected outputs, and the proposed step-by-step solution (subgoals).

Your role is to review the task from the perspective of a professional who has been asked to complete it. Determine whether the task is realistic, whether it can be successfully completed using the provided files and resources, and whether the proposed solution reflects how you would approach the work. Identify any issues, validate the proposed solution, and provide feedback that will help improve both the task quality and the overall pipeline. **You are not expected to create new tasks, rewrite task instructions, or directly edit the task itself—your role is to evaluate the task and provide expert feedback.**



What Makes a Good Task?

What Makes a Good Task?

A strong task represents a realistic professional workflow, requires meaningful end-to-end work, and produces a result that can be clearly verified. These three requirements:

representativeness, complexity, and verifiability — determine which tasks are accepted and which are rejected.

1. Representativeness

The workflow should match real professional practice: the software, files, constraints, and decisions a domain expert would actually encounter. It should feel like work a real professional might delegate to an AI agent, using domain-appropriate tools (a design task should use SolidWorks, Rhino, FreeCAD, Blender, or similar — not an unrelated tool). Strong tasks include realistic inputs — project files, source documents, prior revisions, client notes, vendor data — and may include plausible distractors like outdated versions or conflicting drafts. Avoid toy examples, generic computer actions, or workflows invented only to exercise a UI feature.

2. Complexity

A task should be an end-to-end deliverable, not a single action. Applying a filter, renaming a file, or changing one setting is too narrow — no planning, cross-checking, or judgment required. A strong task requires several interdependent subgoals: inspecting sources, choosing inputs, reconciling constraints, modifying state, validating, and delivering.

- *Too narrow*: "Apply a color filter in DaVinci." A single local edit.
- *Better*: "Move a running cheetah into another race video." Tracking, rotoscoping, compositing, masking, and color matching must all succeed together.

Rule of thumb: the task should take an expert substantial focused time, not a few UI operations.

3. Verifiability

The output must be checkable — ideally deterministically (comparing files, checking numeric values, inspecting geometry, validating formulas, reading backend state). Where exact matching is impossible, evaluation should still reduce to a concrete artifact with an unambiguous rubric. Avoid tasks where success depends on subjective taste or open-ended interpretation.

- *Too open-ended*: "Design an RPG game with monsters." Many wildly different outputs could satisfy it.
- *Better*: "Reproduce the game `mota.exe` using RPGMaker XP." Map geometry, character attributes, and scripted events can be compared against a reference.

A task may allow multiple valid solutions, but it must define which properties count as correct and how an evaluator would accept valid alternatives while rejecting near-misses.



Evaluate the Task

Evaluating the Task

Each task is presented as a complete review package. It may include the assigned professional domain, task instructions, proposed rubric and weights, expected applications and tools, required files, distractors, and a proposed Reference Solution Path. Your role is to examine these materials from the perspective of a professional in the listed field and determine whether the task is realistic, coherent, and completable.

You are **not** expected to perform the task, calculate every answer, produce the requested deliverable, or submit work through the listed applications. Instead, review the materials closely enough to determine whether a competent professional could complete the task using the provided instructions, files, applications, and resources, and whether the proposed solution reflects a reasonable professional approach.

Work through the task in the following order:

1. **Confirm the domain matches your expertise.**

Begin by reviewing the listed sector and occupation. Confirm that the task falls within an area you are qualified to evaluate and that the work described is realistic for a professional in that field. If the task is outside your expertise or appears to belong to a different occupation, skip the task, and let us know on Slack so that we can check that you are mapped to the right domain.

2. **Read and evaluate the *Task Instructions*.**

Read the instructions as though they were sent to you by a colleague. Make sure you understand the objective, the requested deliverable, the relevant context, and any stated constraints. Consider whether a professional would know what to do, what to produce, and how to recognize successful completion. Note any ambiguity, missing context, contradiction, or requirement that would force the professional to make unsupported assumptions.

3. **Review the Rubric and Weights.**

Treat the rubric as a proposed evaluation of successful task completion. Consider whether the rubric categories reflect the most important parts of the work, whether any categories are missing or unnecessary, and whether the assigned weights match their professional importance. Also determine whether the category descriptions are clear, specific, and aligned with what the Task Instructions actually ask the agent to accomplish.

4. **Review the applications, tools, and distractors.**

Determine whether the listed applications and tools are appropriate and sufficient for completing the task. Consider whether any required application is missing or whether an application has been included without a meaningful role.

5. **Examine the supporting files and resources.**

Review the *Required Files* section and open each provided file to understand what it contains, how it is expected to be used, and whether it is usable for the task. Confirm that the files described in the task are present, accessible, and appropriate for producing the requested deliverable. You do not need to complete the analysis contained in the files, but you should inspect them closely enough to determine whether they provide the necessary information. Then review any listed distractors or red herrings and determine whether they create fair difficulty without making the correct solution unavailable, misleading, or ambiguously scored.

6. **Review the Reference Solution Path.**

Read the proposed step-by-step solution from beginning to end. Do not carry out the steps yourself. Instead, determine whether following the proposed path with the listed files, tools, and instructions would reasonably lead to the correct and complete deliverable. Consider whether the sequence reflects how a professional would approach the work, whether any essential steps are missing, and whether any steps are unnecessary, incorrect, or out of order.

7. **Provide your overall expert assessment.**

After reviewing the full package, answer the review questions based on your professional judgment. Identify the most likely failure modes, estimate how long the task itself would take a competent professional, assess its difficulty, and rate the quality of the Task Instructions separately from the quality of the overall task package.

Throughout the review, remember that your responsibility is to **evaluate, not execute or repair**. You should not attempt to complete the underlying task, create a new task, rewrite the *Task Instructions*, modify the files, change the rubric, or edit the proposed solution. When you identify an issue, describe it clearly and specifically so that the task can be improved by the appropriate team. For example, “The task is vague” is less useful than “The instructions do not specify which file should be used to calculate the final rate.”

 Complete your Evaluation
(Submission)

1 Task Instructions

This guide walks you through each question you'll encounter on the platform, with a clear explanation of what each one is asking and how to answer it well. The questions follow the natural order of a review: first the task's setup (domain, applications, and files), then its design (instructions, subgoals, and feasibility), then its evaluation (the golden solution and verifier rubric), and finally your expert judgment (failure modes, quality rating, and time estimates) and submission details.

Keep this guide handy as you work — if you're ever unsure what a question means or how to answer it, refer back to the explanation here.

Domain Alignment

Are the Task Instructions realistic for a professional in the listed domain?

Every task is tagged with an occupational domain (for example, medical, construction, or engineering). This question asks you to judge whether the work in the task is something a real professional in that field would plausibly do as part of their job. Read the task as if you were someone working in that domain: does this feel like a genuine request you might receive or delegate? A task fails this check if the work belongs to a different field entirely, or if it's so artificial that no professional would recognize it as real work. If you answer "No," explain the mismatch and, where possible, point to the domain that actually fits or the changes that would make the task authentic.

Goal & Context Clarity

Do the Task Instructions clearly state the goal and enough context to begin, without requiring unstated assumptions?

This question asks whether a professional could understand what they are being asked to accomplish after reading only the **Task Instructions**. The instructions should clearly explain the objective, provide enough business or professional context to understand the request, and identify any important constraints or background needed to begin the work. Focus only on the written instructions—not the supporting files or other task components. A task fails this check if a professional would need to make significant assumptions or ask follow-up questions before they could confidently begin. If you answer “**No**,” identify the missing context or explain what additional information would be needed.

Procedure in Task Instructions

Do the Task Instructions prescribe a step-by-step procedure?

Professional work requests typically describe **what** needs to be accomplished rather than **how** to accomplish it. This question asks whether the Task Instructions appropriately define the desired outcome without unnecessarily prescribing every action the professional should take. Some procedural guidance may be appropriate when required by regulations, policies, or submission requirements, but the instructions should still allow a professional to apply their own expertise and workflow. A task fails this check if it reads like a detailed checklist of actions instead of a realistic professional request. If you answer “**Yes**,” identify the procedural language that should be removed or rewritten so the instructions state the goal and context only (not how to complete the task).

Correct Deliverable Definition

Based on the Task Instructions alone, is it clear what a completed deliverable would be for this task?

This question asks whether the **Task Instructions** clearly define what the professional is expected to produce. After reading the instructions, a professional should understand what the final deliverable is, when the work is considered complete, and what successful completion looks like without relying on the supporting files, rubric, or proposed solution. A task fails this check if the requested output is vague, incomplete, contradictory, or open to multiple reasonable interpretations. If you answer “**No**,” explain what aspect of the expected deliverable is unclear or missing.

Instruction Quality

Rate the quality of the task instructions as though you received this request from a colleague. (Refer to the [Rating Task Quality](#) in guidelines)

2 Rubrics and Weights

Section Goal: Verify that the golden solution and evaluation rubric accurately assess task completion.

Rubric Weightings

The rubric can be found in the Rubric and Weights table within the task. It lists rubric categories with descriptions and % weights. Do you agree with the relative weightings presented there? That is, in your assessment, do these relative weightings correctly reflect the relative importance of each category to the task?

Unnecessary and Missing Rubric Categories

Are any listed rubric categories unnecessary for judging successful completion, or are any important parts of successful task completion missing from the rubric?

A strong rubric evaluates every critical part of the task without rewarding work that is unrelated or unnecessary. Review each rubric category and determine whether it represents an important aspect of successful task completion. Then consider whether any major requirement described in the Task Instructions is not represented by the rubric. A rubric fails this check if it omits an essential part of the task, includes categories that do not meaningfully contribute to evaluating success, or overemphasizes work that is incidental to the deliverable. Answer “Yes” if you believe categories should be added or removed.

Rubric Category Structure

Should any rubric categories be split or merged so that scoring better matches how the work is actually done?

Sometimes a single rubric category combines multiple independent skills that should be evaluated separately. In other cases, several categories measure nearly the same thing and could reasonably be combined. Review the organization of the rubric and determine whether its structure reflects the natural workflow and major components of the task. A rubric fails this check if categories are overly broad, significantly overlap one another, or divide work into pieces that are unnecessarily granular. Answer “Yes” if you believe reorganizing the categories would produce a clearer or more meaningful evaluation.

Rubric Category Descriptions

Are the rubric category descriptions clear and specific enough to understand what successful completion looks like for each category?

Each rubric category should clearly describe what is being evaluated and what constitutes successful completion. Review the descriptions and determine whether they are specific enough that a reviewer could consistently interpret and apply them without guessing. A rubric fails this check if category descriptions are vague, ambiguous, internally inconsistent, or omit important

details needed to understand what is being scored. If you answer **“No,”** explain which descriptions are unclear and why.

Rubric ↔ Task Instructions Alignment

Does the set of rubric categories match what the Task Instructions ask the agent to produce or accomplish?

The rubric should evaluate the work requested in the Task Instructions—and only that work. Compare the rubric against the Task Instructions and determine whether every major requirement described in the instructions is reflected in the evaluation. Likewise, confirm that the rubric does not assess work that the Task Instructions never ask the agent to perform. A rubric fails this check if important task requirements are missing from the evaluation, or if the rubric scores outputs, behaviors, or information that are unrelated to the requested deliverable. If you answer **“No,”** explain where the mismatch occurs and whether the Task Instructions or the rubric should be updated.

3 Resources for Completion

Applications & Tools

Are the listed applications suitable and sufficient for completing the deliverable? Answer "No" if any necessary applications are missing from the list. Following this, for each application you will assess whether it is appropriate and required for the task.

Each task lists the applications the agent is expected to use. This question asks two things at once: are the listed applications the *right* ones for this work (appropriate), and can the task actually be completed using only what's listed (sufficient)? A task fails this check if it lists an application that has no role in the work, or if completing the task would require software that isn't listed.

Required Files

Open each of the required files, which are attached to the task via a .ZIP folder, and review them carefully. In this section, you will use your industry experience to assess whether each file is applicable and appropriate to the task described.

First, confirm whether you are able to open all of the required files. Then, assess each file individually via the questions below.

Open each file in the provided .ZIP folder and review it closely enough to understand its purpose, structure, and relevance. Consider whether the file formats, contents, level of detail, terminology, and organization are consistent with real industry work. Also confirm that each file is applicable to the task and supports the type of analysis or deliverable requested in the Task Instructions. A task fails this check if the files appear artificial, contain implausible or irrelevant information, do not match their stated purpose, or would not reasonably be used by a professional performing this work. If you answer “**No**,” specify which files are not realistic and applicable to the task described. Suggest improvements to the content of the files to make them more realistic as applicable.

Deliverable Feasibility

Could a professional in this field complete the requested deliverable correctly using only the provided task instructions, files, and applications (including known distractors)?

This question asks you to evaluate the task as a complete package. Consider the Task Instructions, supporting files, applications, and known distractors together and determine whether they provide everything needed for a competent professional to produce the requested deliverable. You are **not** expected to complete the task yourself. Instead, determine whether all required information, resources, and tools are present and whether there are any missing dependencies, conflicting requirements, or impossible steps that would prevent successful completion. A task fails this check if the requested deliverable cannot reasonably be produced using only the provided materials. If you answer “**No**,” explain what prevents successful completion.

Reference Solution Path

If you followed the proposed Reference Solution Path with the listed Task Instructions, Apps Used, and Required Files, could you reach the correct, complete deliverable?

The Reference Solution Path represents the proposed workflow for completing the task. This question asks you to evaluate whether those steps form a realistic and complete path to the expected deliverable. You are **not** expected to execute the solution or verify every calculation. Instead, review the sequence of subgoals and determine whether they are logically ordered, use the appropriate files and applications, and reflect how a professional would reasonably approach the work. A solution path fails this check if it omits essential steps, relies on unavailable information, contains unnecessary or incorrect actions, or would not reasonably produce the requested deliverable. If you answer “**No**,” explain where the solution path breaks down and what changes would make it accurate and complete.

4 Overall Assessment

Task Validity

After assessing the task, would you consider the task valid or invalid?

A binary choice. After reviewing the entire task, please consider if this task is valid or invalid.

A **valid task** is one that a competent professional could reasonably complete using the provided Task Instructions, files, applications, and other supporting materials. The task has a clear objective, a defined deliverable, appropriate resources, and no major issues that would prevent successful completion. Minor imperfections or opportunities for improvement do **not** make a task invalid.

An **invalid task** contains one or more critical issues that prevent the task from being completed or fairly evaluated. Examples include missing or unusable required files, contradictory or impossible instructions, an undefined or unattainable deliverable, missing required applications or resources, or other major flaws that would prevent a competent professional from successfully completing the work. If a task is invalid, identify the specific issue that makes completion impossible or unreasonable.

Estimated Time to Complete (ETC)

Approximately how long would it take a competent professional in this field to complete this task from start to completion?

Estimate, in minutes, how long the task would take a skilled human professional, someone who knows the domain and the software, working efficiently without prior knowledge of this specific task. Count the full effort from reading the instructions to producing the final deliverable, including gathering information from the provided materials, but not detours, mistakes, or learning time a competent professional wouldn't need. This doesn't have to be precise to the minute; a realistic estimate is enough. ETC gives us a consistent measure of task size and depth, and helps confirm the task is the substantial, long-form work we're aiming for rather than a quick one-step job.

Estimated Difficulty

How would a mid-level professional in the field rate the difficulty of the task assuming no access to LLMs?

Please [see this detailed guide for more information](#).

5 Complete the Submission

Time to Complete (minutes)

How long did it take you to complete this submission?

Unlike the ETC question, which asks you to estimate how long the task would take a professional, this question asks about you: how many minutes did you actually spend on this review, from opening the task to submitting your answers? Count all of it: reading the instructions, opening and checking the files, working through each question, and writing your explanations. Report your real time honestly, even if it feels long or short; there's no right answer. This data helps us understand the true cost of reviewing, spot tasks that are unusually burdensome, and plan reviewer workloads realistically.

AI Use Attestation

The use of AI or Large Language Models (LLMs) is strictly prohibited during this review. Please check the box below to confirm that you completed this task entirely without AI assistance.

This is a formal confirmation, not a question with a judgment to make. The entire value of this review rests on independent human expertise. If you cannot truthfully check the box, do not submit the review.



Rating Task Instructions

Rating Task Difficulty

Imagine assigning this task to a competent **mid-level professional** working in the relevant field. Assume they have the appropriate software, files, and typical workplace resources available, **but no access to LLMs or AI assistants**.

Rate the difficulty of completing the task itself—not the quality of the instructions.

Consider questions such as:

- How much domain expertise is required?
- How much independent reasoning or judgment is needed?
- How technically or procedurally complex is the work?
- How many meaningful steps or decisions are involved?
- Would this be considered routine work for a professional, or something unusually challenging?

Focus on the work required to successfully complete the task, not how long it takes.

Score	Description
5 – Very Difficult	The task requires deep domain expertise, advanced technical knowledge, or substantial independent judgment. A mid-level professional would find it challenging and would need to solve complex problems, synthesize information from multiple sources, or make numerous consequential decisions. This is work that would typically be assigned to experienced professionals or specialists.
4 – Difficult	The task requires solid professional knowledge and involves multiple non-trivial decisions or technical steps. A mid-level professional should be able to complete it successfully but would need to think carefully, apply expertise throughout the process, and avoid several potential pitfalls.
3 – Average Difficulty	The task reflects routine professional work. It requires relevant domain knowledge and some judgment, but the workflow is generally familiar to

	someone with experience in the field. Most mid-level professionals could complete it without unusual difficulty.
2 – Easy	The task is straightforward and follows well-established procedures. It requires only limited professional judgment and consists primarily of routine actions that someone in the field performs regularly. Most mid-level professionals would complete it with little effort.
1 – Very Easy	The task is extremely simple or administrative in nature. It requires little specialized knowledge, minimal decision-making, and almost no technical complexity. A professional could complete it quickly with little cognitive effort.



UNDER CONSTRUCTION: REVIEWER GUIDELINES

The following items are known as “Error Flags”. These are what reviewers mark tasks with in order to signal that the tasks do not meet specifications for a certain reason.

- **Low-Detail / Surface-Level Analysis:** Vague or generic notes that do not engage the domain, task-specific instructions, apps, files/distractors, rubric categories, or reference solution path.
- **Instruction or Deliverable Misread:** Misstates the goal, whether instructions are procedural, or what “done” means (wrong artifact type, missing Submit/end state, or invents requirements not in the Task Instructions).
- **Rubric Misjudgment:** Vastly incorrect weight of Agree/Higher/Lower calls, skipped categories, or wrong claims about missing/unnecessary categories or rubric↔instruction alignment relative to the Rubric and Weights table.
- **Resources / Feasibility Misassessment:** Incorrect claim that apps, required files, distractors, or the reference solution path make the task unfair, unusable, or unsolvable.
- **Logical Inconsistency:** Final rating contradicts the written conditionals or each other (e.g., “deliverable is clear” paired with notes describing an unclear end state).
- **Spam, Low-Quality, & AI Output:** The response is vague, generic, or obviously automated. It lacks specific asset references (e.g., direct quotes, category/file/path names) or contains clear AI hallucinations and features not actually present in the asset.